MD. SAIF ALI

Full-Stack Software Engineer | Node.js | MERN Stack | AI / RAG
Durgapur, West Bengal, India | +91 78108 71384 | saifmd9536@gmail.com
linkedin.com/in/mdsaifali09 | github.com/mdsaifalidev | buildwithsaif.vercel.app

SUMMARY

Full-stack engineer with 3+ years of experience on Node.js backends and event-driven microservices. Now building RAG and MCP systems at HashCash Consultants. Also builds REST and GraphQL APIs and React frontends, and ships containerized services to AWS via CI/CD.

SKILLS

Languages: JavaScript (ES6+), TypeScript, Python, SQL
Generative AI: RAG, MCP, LangChain, OpenAI & Gemini LLMs, Semantic Search, Prompt Engineering
Backend: Node.js, Express.js, REST APIs, GraphQL, Microservices, RabbitMQ (Event-Driven & RPC), JWT, OAuth2, Stripe, PayPal
Frontend: React.js, Next.js, TanStack Query, Zustand, Tailwind CSS, Socket.IO (WebSockets)
Databases: PostgreSQL, MongoDB, Redis, Prisma ORM, Vector Database (pgvector)
Cloud & DevOps: Docker, Kubernetes, AWS (EC2, S3, IAM, Route 53, CloudWatch, ECR), GitHub Actions (CI/CD), Git
Testing & Process: Jest, Mocha, Chai, Agile/Scrum

EXPERIENCE

Software Engineer | HashCash Consultants, Kolkata, West Bengal | **Sep 2025 - Present**
- Designed a multi-tenant Retrieval-Augmented Generation (RAG) chatbot platform (Node.js, Python, LangChain, OpenAI, Gemini, pgvector) in which every tenant has its own isolated knowledge base. Manual information retrieval dropped by ~30%.
- Built the RAG pipeline: document ingestion, chunking, embeddings, vector storage and semantic search.
- Built a read-only Model Context Protocol (MCP) server. AI agents can search and retrieve organizational knowledge through it but cannot change data or reach raw systems.
- Designed and shipped a GraphQL API (schemas, resolvers, query/mutation contracts) that cut payload size and round trips.
- Cut p95 REST API latency by 20% in a production MERN app by reworking database indexes and query logic.

Full Stack Developer | Ayanlabs Solutions Pvt. Ltd., Remote | **Aug 2024 - Aug 2025**
- Built a Node.js/Express.js microservices backend with 6+ services. Packaged them in Docker (images in AWS ECR) and deployed to Kubernetes on AWS EKS with Helm charts.
- Deployed applications on AWS EC2 and kept images, PDFs, videos and backups in S3. Managed IAM users, roles and permissions, pointed domains through Route 53, and watched servers and apps in CloudWatch logs.
- Connected the services with asynchronous RabbitMQ messaging (RPC and event-driven) so core business logic stayed decoupled.
- Set up GitHub Actions CI/CD with Mocha/Chai unit and integration tests. Manual release effort fell ~60%; test coverage reached 80%+.

Full Stack Developer Intern | Sasefied, Remote | **Jan 2024 - Jun 2024**
- Built features for a PERN-stack app (PostgreSQL, Express.js, React.js, Node.js) and took them through production deployment, including Stripe payments and RabbitMQ order confirmations.
- Split high-traffic workloads into microservices. PostgreSQL tuning, Redis caching and TanStack Query made page loads ~40% faster.

Full Stack Developer Intern | E-SUTRA Technologies Pvt. Ltd., Remote | **Jun 2023 - Nov 2023**
- Built MERN-stack features, including REST APIs and payment gateway integrations.
- Added OAuth2 social login and iterated on the UI/UX with an Agile/Scrum team.

PROJECTS

InsightFlow | Next.js, Fastify, FastAPI, Polars, PostgreSQL, Redis, S3, Docker | useinsightflow.vercel.app | GitHub
- Built an app that turns an uploaded CSV or Excel file (up to 100 MB) into an interactive dashboard and an AI-written executive brief. A 1M-row, 86.7 MB file takes ~12 s; filtered queries run at p95 368-377 ms.
- The model never writes a number. A Polars engine computes every figure, Gemini refers to them by fact ID, and a validator rejects anything not backed by a computed fact.
- Split into a Fastify auth service (RS256 JWT/JWKS), a FastAPI data service and an ARQ worker in a Turborepo monorepo, with SSE streaming job progress to the browser. Has 189 unit and integration tests plus Playwright E2E.

Portfolio AI Assistant | Next.js, Vercel AI SDK, Gemini, Upstash Redis, TypeScript | buildwithsaif.vercel.app | GitHub
- Built a RAG assistant for my portfolio site that answers from the site's own content, streams replies and links each [n] citation to its source section. Retrieval runs in memory over an index built at build time, so there is no vector database.
- Added prompt-injection guardrails and Upstash rate limits (per visitor and site-wide). A 24-question eval script passes 24/24, and the first token arrives in ~1.2 s.

EDUCATION

Bachelor of Computer Applications (BCA) | Bengal College of Engineering and Technology (BCET), Durgapur | **2022 - 2025**

CERTIFICATIONS

NodeJS - The Complete Guide (MVC, REST APIs, GraphQL, Deno) | Udemy
Certificate: ude.my/UC-d9afe7b1-4a9f-4111-8355-cd1d05775d53